

Docker端口&容器互联

@姜桥

目录

- 外部容器访问
- Docker网络模型
- Docker网络模式
- Pipework
- Mobike实践

外部容器访问

例如:Docker中安装一个Mysql:

- 1、从Docker镜像仓库获取Mysql官方镜像包

```
docker pull docker.io/mysql
```

- 2、启动Mysql容器并指定端口映射


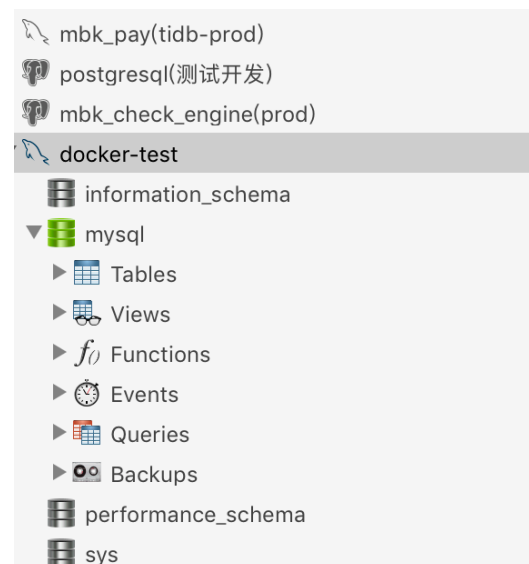
//创建宿主挂载目录:

```
sudo mkdir -p /data/mysql/logs /data/mysql/conf /data/mysql/data -p
```

```
docker run --name mysql -p 3306:3306 -v /data/mysql/data:/var/lib/mysql -v /data/mysql/conf:/etc/mysql/conf.d  
-e MYSQL_ROOT_PASSWORD=123456 -d mysql
```

```
[jiangqiao@localhost logs]$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
b3b8139754c7	mysql	"docker-entrypoint.s..."	34 minutes ago	Up 34 minutes	0.0.0.0:3306->3306/tcp, 33060/tcp	mysql



通过宿主机IP访问容器中的MySQL

外部容器访问

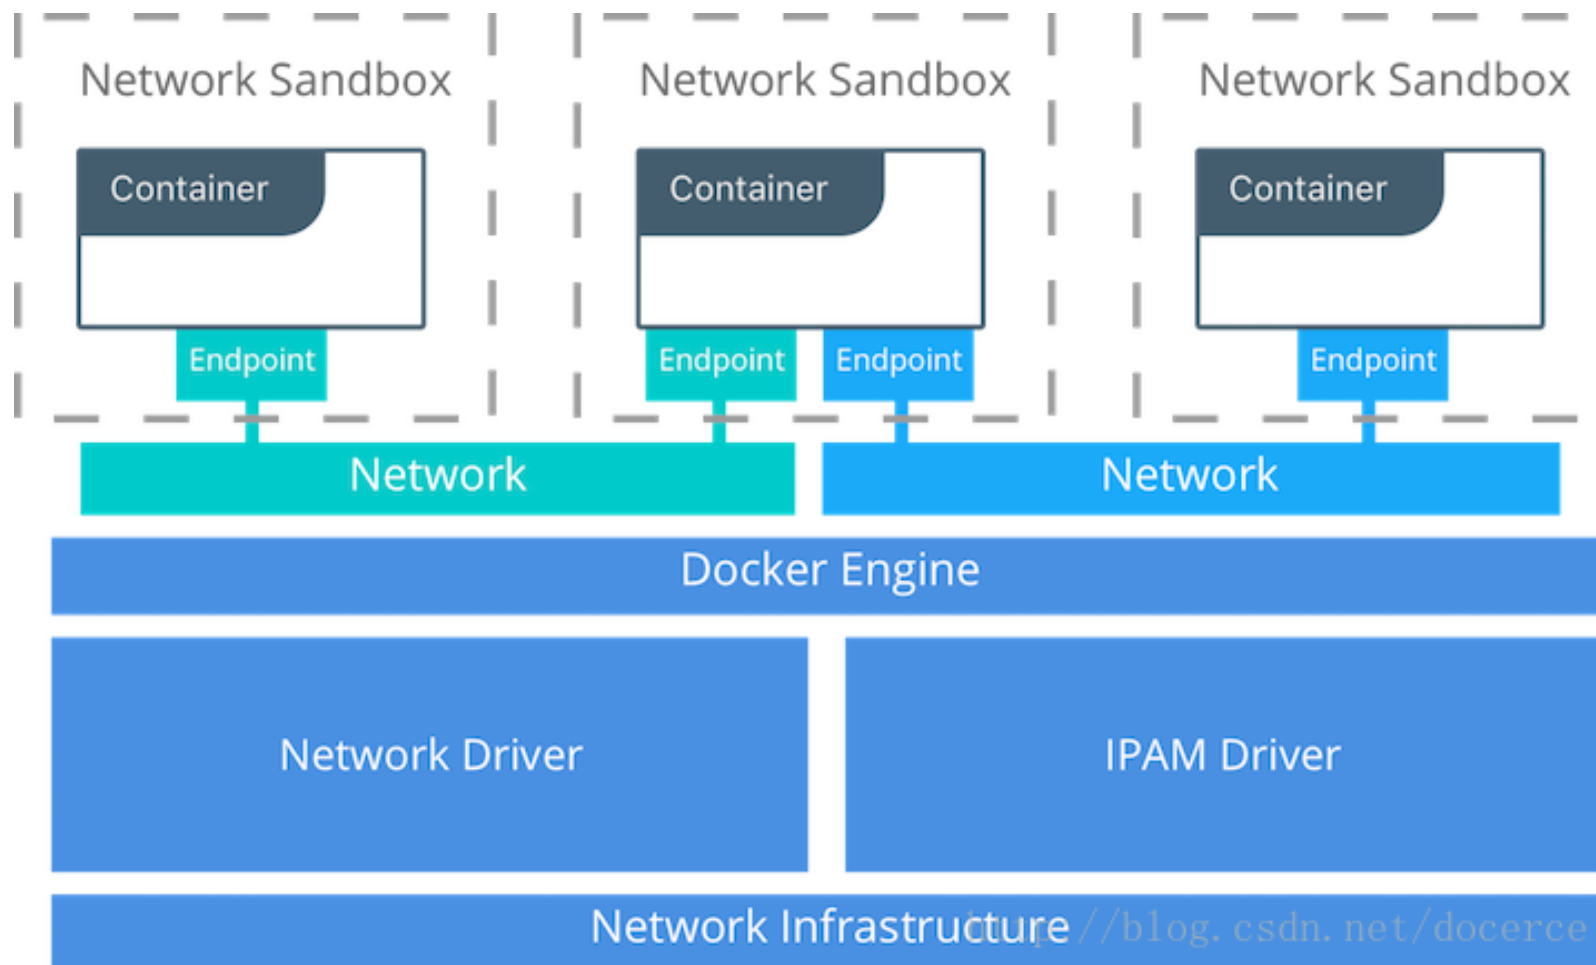
Docker容器中运行的应用，要让外部访问，可以通过**-P 或-p** 参数来指定端口映射。

在上面的例子中就是绑定了宿主机的3306端口到容器的3306端口。

可以通过 `docker port`来查看当前映射的端口配置和绑定的地址

```
[[jiangqiao@localhost logs]$ docker port b3b8139754c7 3306  
0.0.0.0:3306
```

Docker网络模型



容器网络模型 (CNM)

- CNM的理念是提供可以跨不同网络基础架构、可实现移植的应用。这种模型虚拟了所有操作系统和基础架构的不可知性，所以应用无论在何种基础架构上都可以有一样的体验；

libnetwork

- libnetwork是基于CNM模型将docker engine 和 lib container中网络相关代码抽离出来，合成的一个单独的库；
- 基于CNM模型概念，进一步对Docker的网络结构细分为：network、sandbox、endpoint三层；

Docker网络模型

CNM结构

■ sandbox

sandbox 实现了容器内部的网络栈，它定义了容器的虚拟网卡，路由表和DNS等配置，其实就是一个标准的linux network namespace实现。

■ network

network 是一个抽象的概念，可以理解为一个网络插件，或者是网络的Driver，是由Linux桥接，Vlan等来实现的；例如 Docker中原生的Driver包括单主机的none、host、bridge，joined container 和多主机的overlay、macvlan，第三方Driver包括多主机的 flannel、weave、calico 等。

网络上收集了所有连接在其上的端点，并实现了这些端点的互连接。

■ endpoint

network实现了一个第三方的网络栈，sanbox则实现了容器内部的网络栈，而它们的互连则是通过endpoint，endpoint实现了veth pair（虚拟网络设备）；一个endpoint就表示一对veth pair，一端挂在容器中，另一端挂在network中。

Docker网络模式

“不同的网络方案如何集成到Docker网络模型中而不改变原有结构？”

Docker提供的网络模式：

单主机网络：

- Bridge
- Host
- Container
- None

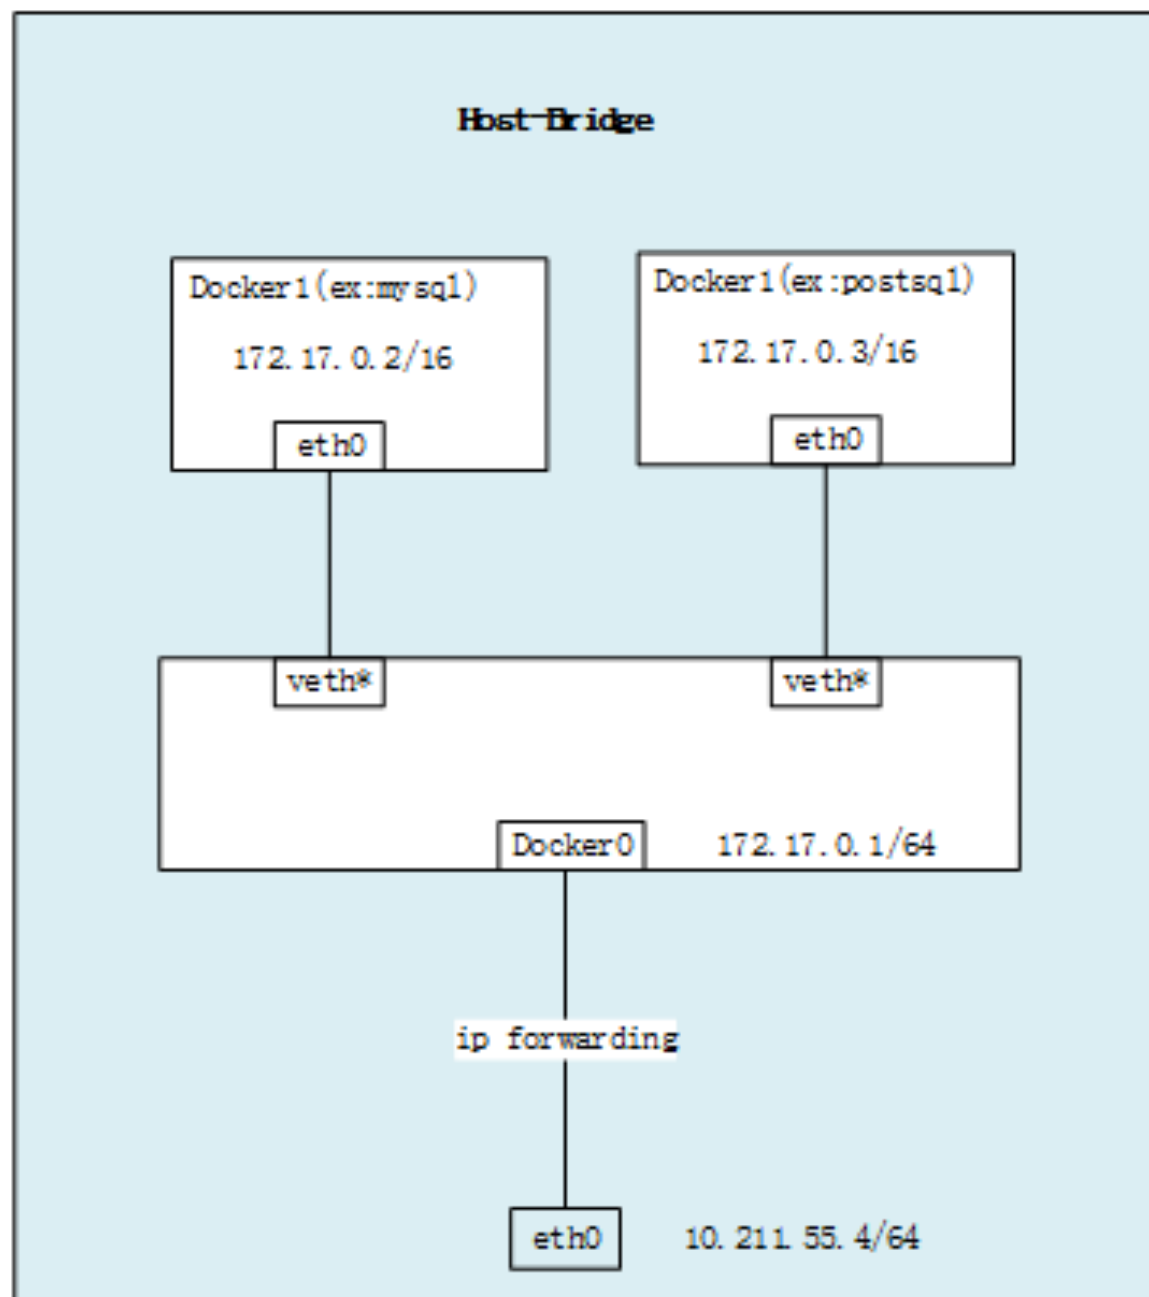
多主机网络

- overlay、macvlan、flannel、weave、calico

网络模式-Bridge

```
docker run --net=bridge --name mysql -p 3306:3306 -v /data/mysql/data:/var/lib/mysql -v /data/mysql/conf:/etc/mysql/conf.d -e MYSQL_ROOT_PASSWORD=123456 -d mysql
```

```
docker run --net=bridge --name postgres1 -e POSTGRES_PASSWORD=123456 -p 54321:5432 -d postgres:9.4
```



- 在Bridge模式下，*docker*进程启动时会在主机上创建一个名为*docker0*的虚拟网桥，此主机上的*docker*容器会连接到这个虚拟网桥上；
- 虚拟网桥的工作方式和物理交换机类似，这样主机上的所有容器就通过交换机连在了一个二层网络中；
- *docker0*网桥从子网中分配一个IP给容器使用，并设置*docker0*的IP地址为容器的默认网关；
- 容器启动后会在主机上创建一对虚拟网卡 *veth pair* 设备，*docker*将*veth pair*设备的一端放在新创建的容器中，并命名为*eth0*作为容器的网卡，另一端放在主机中，以*vethxxx*这样类似的名字命名，并将这个设备加入到*docker0*网桥中；

网络模式-Bridge

```
[jiangqiao@localhost data]$ ifconfig
docker0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:1aff:fec1:b1d8 prefixlen 64 scopeid 0x20<link>
    ether 02:42:1a:c1:b1:d8 txqueuelen 0 (Ethernet)
    RX packets 239 bytes 26531 (25.9 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 387 bytes 31872 (31.1 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.211.55.4 netmask 255.255.255.0 broadcast 10.211.55.255
    inet6 fe80::21c:42ff:fea4:e10a prefixlen 64 scopeid 0x20<link>
    inet6 fdb2:2c26:f4e4:0:21c:42ff:fea4:e10a prefixlen 64 scopeid 0x0<global>
    ether 00:1c:42:a4:e1:0a txqueuelen 1000 (Ethernet)
    RX packets 370649 bytes 528536845 (504.0 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 125798 bytes 7728069 (7.3 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 92 bytes 7776 (7.5 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 92 bytes 7776 (7.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

veth38f43dc: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::a83c:3dff:fe70:9ab0 prefixlen 64 scopeid 0x20<link>
    ether aa:3c:3d:70:9a:b0 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 9 bytes 729 (729.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

veth71b8b5b: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::cda:6aff:fea2:737b prefixlen 64 scopeid 0x20<link>
    ether 0e:da:6a:a2:73:7b txqueuelen 0 (Ethernet)
    RX packets 53 bytes 9084 (8.8 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 87 bytes 7242 (7.0 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:82:75:69 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
[jiangqiao@localhost data]$ docker inspect e2807260e87b
```

```
{
  "NetworkSettings": {
    "Bridge": "",
    "SandboxID": "18106dc7ed0c82e21b9f1ec75d7a2d1243a0c3349f6591c45c8d726108c12c57",
    "HairpinMode": false,
    "LinkLocalIPv6Address": "",
    "LinkLocalIPv6PrefixLen": 0,
    "Ports": {
      "5432/tcp": [
        {
          "HostIp": "0.0.0.0",
          "HostPort": "54321"
        }
      ]
    },
    "SandboxKey": "/var/run/docker/netns/18106dc7ed0c",
    "SecondaryIPAddresses": null,
    "SecondaryIPv6Addresses": null,
    "EndpointID": "b157af56181afef7d89e0bc117040bc2741ed88106d4684f96550074c4e797c4",
    "Gateway": "172.17.0.1",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "IPAddress": "172.17.0.3",
    "IPPrefixLen": 16,
    "IPv6Gateway": "",
    "MacAddress": "02:42:ac:11:00:03",
    "Networks": {
      "bridge": {
        "IPAMConfig": null,
        "Links": null,
        "Aliases": null,
        "NetworkID": "35aa8a01e83ae1739d809b1f89263c4bc820d2ca8bf6a8ef2929d3385f17a1fa",
        "EndpointID": "b157af56181afef7d89e0bc117040bc2741ed88106d4684f96550074c4e797c4",
        "Gateway": "172.17.0.1",
        "IPAddress": "172.17.0.3",
        "IPPrefixLen": 16,
        "IPv6Gateway": "",
        "GlobalIPv6Address": "",
        "GlobalIPv6PrefixLen": 0,
        "MacAddress": "02:42:ac:11:00:03",
        "DriverOpts": null
      }
    }
  }
}
```

Docker端口映射

Bridge模式时Docker的默认网络模式，不写--net参数，就是bridge模式。在这种模式下，容器如果需要联网，则需要采用NAT（网络地址端口转换）方式；

(1)、默认情况下，容器可以通过SNAT（出去的时候改变源地址，回来的时候改变目的地址）的方式主动访问到外部网络的连接。

```
sudo iptables -t nat -nL
```

```
Chain POSTROUTING (policy ACCEPT)
target     prot opt source                destination
MASQUERADE all  --  172.17.0.0/16          0.0.0.0/0
RETURN     all  --  192.168.122.0/24      224.0.0.0/24
RETURN     all  --  192.168.122.0/24      255.255.255.255
MASQUERADE tcp  --  192.168.122.0/24      !192.168.122.0/24    masq ports: 1024-65535
MASQUERADE udp  --  192.168.122.0/24      !192.168.122.0/24    masq ports: 1024-65535
MASQUERADE all  --  192.168.122.0/24      !192.168.122.0/24
POSTROUTING_direct all --  0.0.0.0/0             0.0.0.0/0
POSTROUTING_ZONES_SOURCE all --  0.0.0.0/0             0.0.0.0/0
POSTROUTING_ZONES all --  0.0.0.0/0             0.0.0.0/0
MASQUERADE tcp  --  172.17.0.2            172.17.0.2          tcp dpt:3306
MASQUERADE tcp  --  172.17.0.3            172.17.0.3          tcp dpt:5432
```

上述规则将所有源地址在172.17.0.0/16网端的容器，通过宿主机网卡IP进行外部网络访问。

Docker端口映射

(2)、默认情况下，外部网络无法访问到容器，可以在docker run时通过 `-p`或`-P`参数启动。其实就是在本地的iptables的nat表中添加相应的规则。

```
sudo iptables -t nat -nL
```

```
Chain DOCKER (2 references)
target     prot opt source                destination
RETURN     all  --  0.0.0.0/0              0.0.0.0/0
DNAT        tcp  --  0.0.0.0/0              0.0.0.0/0          tcp dpt:3306 to:172.17.0.2:3306
DNAT        tcp  --  0.0.0.0/0              0.0.0.0/0          tcp dpt:54321 to:172.17.0.3:5432
```

这里的规则映射了`0.0.0.0`，意味着将接收主机来自所有接口的流量。可以通过 `-p IP:host_port:container` 或者 `-p IP::port` 来指定允许访问容器的主机上的IP、端口等，以指定更严格的规则；

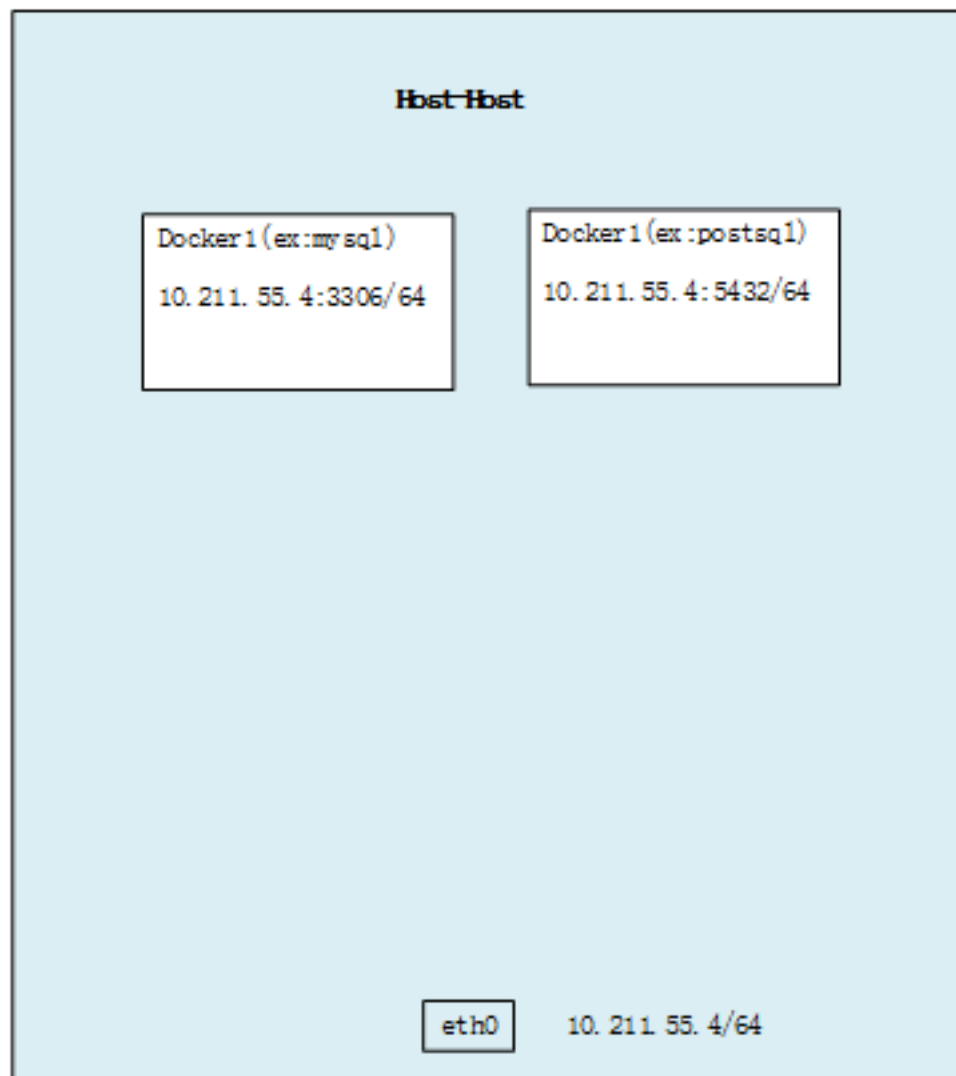
```
[[jiangqiao@localhost data]$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
e2807260e87b	postgres:9.4	"docker-entrypoint.s..."	4 hours ago	Up 4 hours	0.0.0.0:54321->5432/tcp	postgres1
f94c2b916597	mysql	"docker-entrypoint.s..."	5 hours ago	Up 5 hours	0.0.0.0:3306->3306/tcp, 33060/tcp	mysql

网络模式-Host

```
docker run --net=host --name mysql -p 3306:3306 -v /data/mysql/data:/var/lib/mysql -v /data/mysql/conf:/etc/mysql/conf.d  
-e MYSQL_ROOT_PASSWORD=123456 -d mysql
```

```
docker run --net=host --name postgres1 -e POSTGRES_PASSWORD=123456 -p 54321:5432 -d  
postgres:9.4
```



- 在*Host*模式下，容器将不会获得一个独立的*Network Namespace*,而是和宿主机共用一个*Network Namespace*;
- 容器也不会虚拟出自己的网卡，配置自己的*IP*，而是使用宿主机的*IP*和端口；
- 但是，容器的其他方面，如文件系统、进程列表等还是通过相应的*namespace*与宿主机隔离的；

网络模式-Host

```
[[jiangqiao@localhost data]$ ifconfig
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:1aff:fec1:b1d8 prefixlen 64 scopeid 0x20<link>
    ether 02:42:1a:c1:b1:d8 txqueuelen 0 (Ethernet)
    RX packets 245 bytes 26795 (26.1 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 396 bytes 32463 (31.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.211.55.4 netmask 255.255.255.0 broadcast 10.211.55.255
    inet6 fe80::21c:42ff:fea4:e10a prefixlen 64 scopeid 0x20<link>
    inet6 fdb2:2c26:f4e4:0:21c:42ff:fea4:e10a prefixlen 64 scopeid 0x0<global>
    ether 00:1c:42:a4:e1:0a txqueuelen 1000 (Ethernet)
    RX packets 372579 bytes 528706119 (504.2 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 126828 bytes 7888064 (7.5 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 209 bytes 45794 (44.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 209 bytes 45794 (44.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

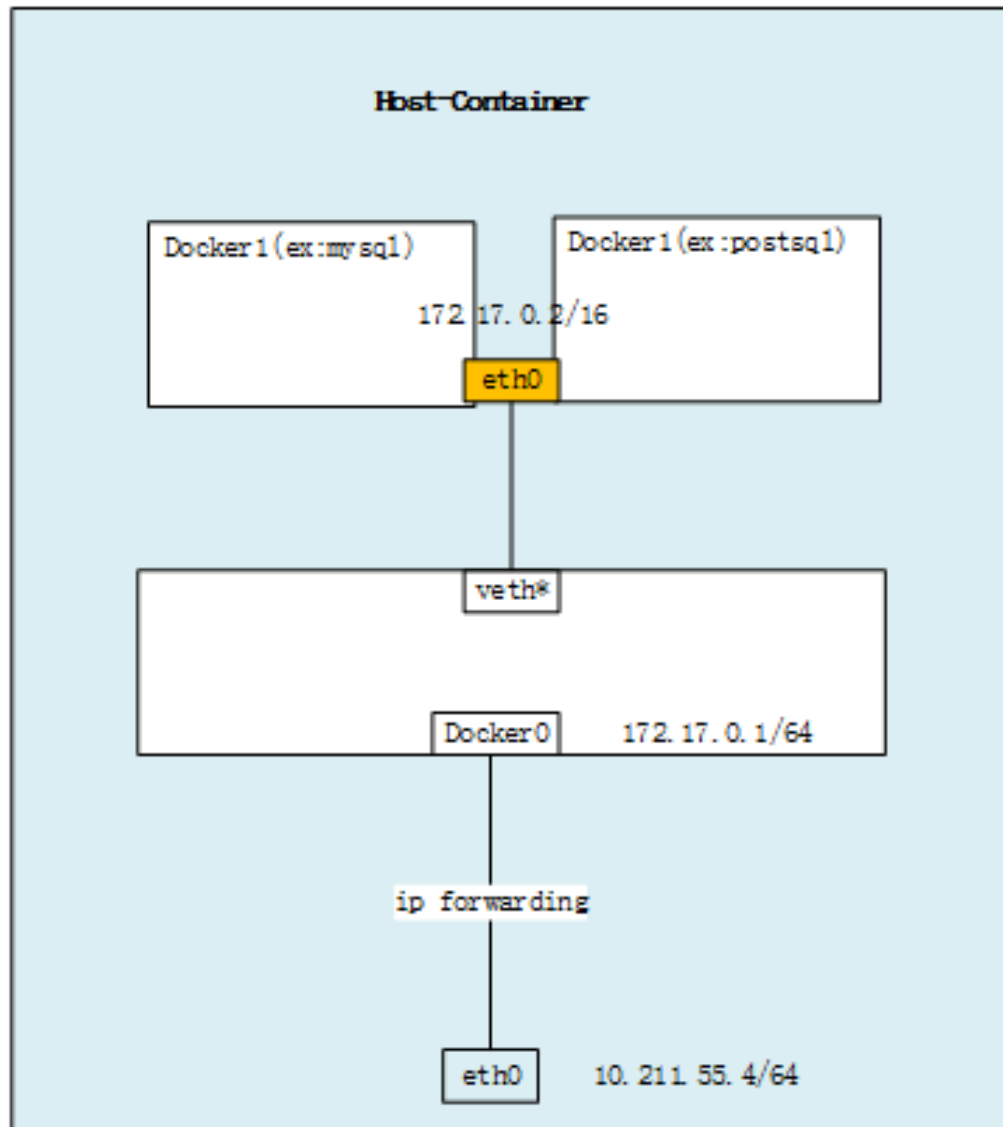
virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:82:75:69 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "fba17826a9d967d5a7d814210bdd1034af23fac1cd6b4bdeffddcc6404dc2736",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/default",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {
    "host": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "4f5bf4f7a37beb738ad487357c77634a4881f1714fea128b2af6fab9239bb330",
      "EndpointID": "2a7ed80de47b29c58433aaeaa0a34778f6c7c363099afac5a6f8b0688dfb0",
      "Gateway": "",
      "IPAddress": "",
      "IPPrefixLen": 0,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "",
      "DriverOpts": null
    }
  }
}
```

网络模式-Container

```
docker run --net=bridge --name mysql -p 3306:3306 -p 5432:5432 -v /data/mysql/data:/var/lib/mysql -v /data/mysql/conf:/etc/mysql/conf.d -e MYSQL_ROOT_PASSWORD=123456 -d mysql
```

```
docker run --net=container:mysql --name postgres1 -e POSTGRES_PASSWORD=123456 -d postgres:9.4
```



- 在 *Container* 模式下，可以指定新创建的容器和已经存在的一个容器共享一个 *Network Namespace*，而不是和宿主机共享；
- 新创建的容器不会创建自己的网卡、配置自己的 *IP*，而是和一个指定的容器共享 *IP*、端口范围；
- 两个容器除了网络方面，其他的如文件系统、进程列表等还是通过不同的 *namespace* 进行隔离；
- 两个容器的进程可以通过 *lo* 网卡设备进行通信；

网络模式-Container

```
[[root@localhost ~]# ifconfig
docker0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:1eff:fe8a:19ee prefixlen 64 scopeid 0x20<link>
    ether 02:42:1e:8a:19:ee txqueuelen 0 (Ethernet)
    RX packets 25 bytes 2552 (2.4 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 58 bytes 5520 (5.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.211.55.4 netmask 255.255.255.0 broadcast 10.211.55.255
    inet6 fe80::21c:42ff:fea4:e10a prefixlen 64 scopeid 0x20<link>
    inet6 fdb2:2c26:f4e4:0:21c:42ff:fea4:e10a prefixlen 64 scopeid 0x0<gl>
    ether 00:1c:42:a4:e1:0a txqueuelen 1000 (Ethernet)
    RX packets 382667 bytes 529681292 (505.1 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 132426 bytes 11125198 (10.6 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 1498 bytes 475914 (464.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1498 bytes 475914 (464.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

veth410fce2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet6 fe80::bc3c:edff:fe21:c866 prefixlen 64 scopeid 0x20<link>
    ether be:3c:ed:21:c8:66 txqueuelen 0 (Ethernet)
    RX packets 25 bytes 2902 (2.8 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 66 bytes 6168 (6.0 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:82:75:69 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "c6d6e4cc394fb5f502128910b3e182091d87057d6b7d2383e9c018861dfc6631",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {
    "3306/tcp": [
      {
        "HostIp": "0.0.0.0",
        "HostPort": "3306"
      }
    ],
    "33060/tcp": null
  },
  "SandboxKey": "/var/run/docker/netns/c6d6e4cc394f",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "0837a17c6649f3e7b881a39841305256e2c0d8664aa72807abae08c1f451185c",
  "Gateway": "172.17.0.1",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "172.17.0.2",
  "IPPrefixLen": 16,
  "IPv6Gateway": "",
  "MacAddress": "02:42:ac:11:00:02",
  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "31d8f1be80740665ad6e483361b3f0a6e62786147877025e11f281a47fd15f8c",
      "EndpointID": "0837a17c6649f3e7b881a39841305256e2c0d8664aa72807abae08c1f451185c",
      "Gateway": "172.17.0.1",
      "IPAddress": "172.17.0.2",
      "IPPrefixLen": 16,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "02:42:ac:11:00:02",
      "DriverOpts": null
    }
  }
}
```

指定容器

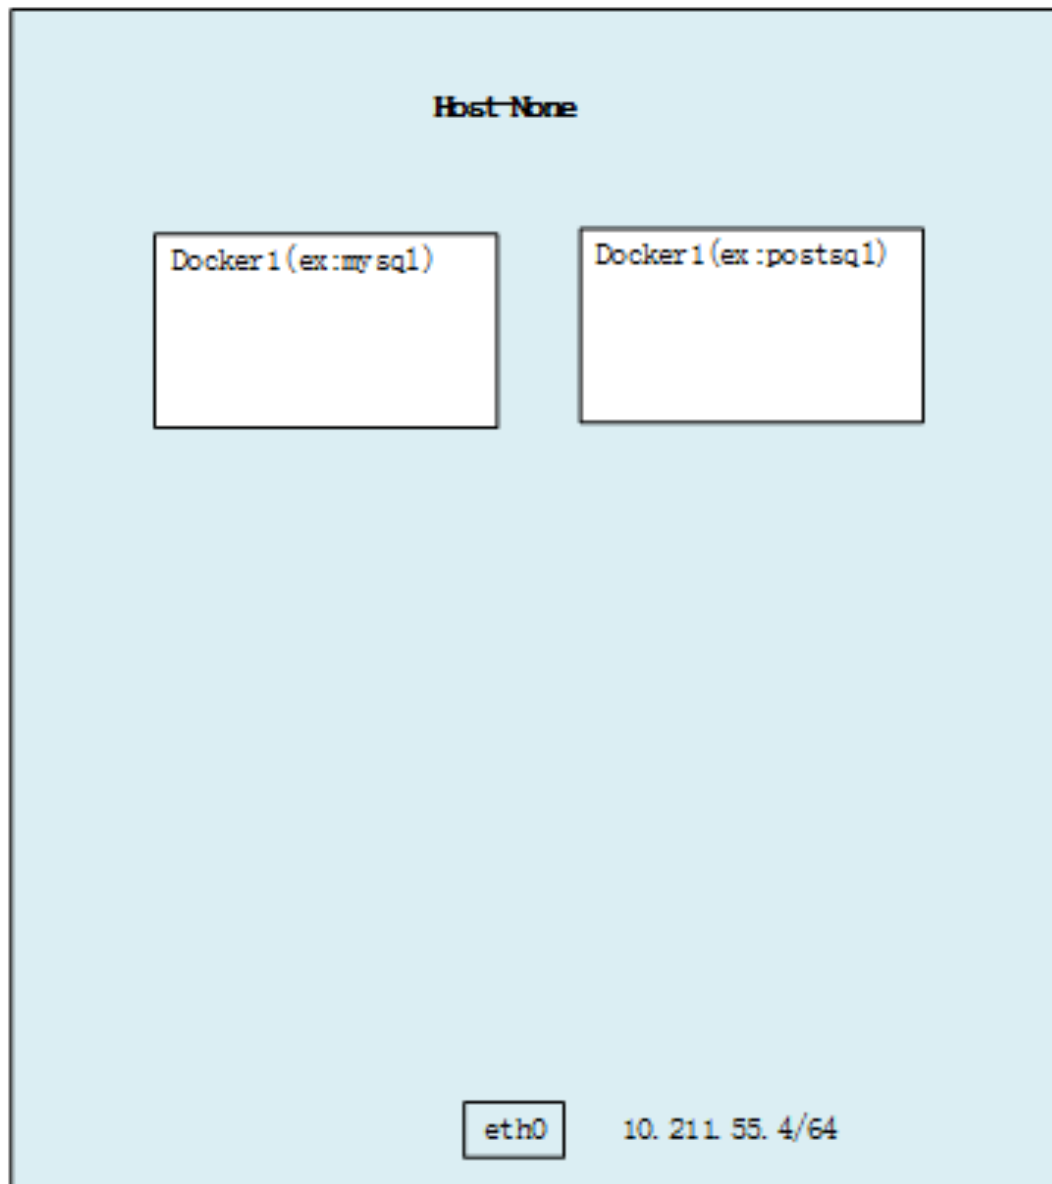
```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {}
}
```

被指定容器

网络模式-None

```
docker run --net=none --name mysql -p 3306:3306 -p 5432:5432 -v /data/mysql/data:/var/lib/mysql -v /data/mysql/conf:/etc/mysql/conf.d -e MYSQL_ROOT_PASSWORD=123456 -d mysql
```

```
docker run --net=none --name postgres1 -e POSTGRES_PASSWORD=123456 -d postgres:9.4
```



- 在None模式下, *Docker*容器拥有自己的*Network Namespace*, 但是, 并不为*Docker*容器进行任何网络配置;
- 这个*Docker*容器没有网卡、*IP*、路由等信息, 需要我们自己为*Docker*容器添加网卡、配置*IP*等;

网络模式-None

Mysql容器

```
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.255
    inet6 fe80::42:f1ff:feef:c35f prefixlen 64 scopeid 0x20<link>
    ether 02:42:f1:ef:c3:5f txqueuelen 0 (Ethernet)
    RX packets 38 bytes 16061 (15.6 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 86 bytes 9139 (8.9 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.211.55.4 netmask 255.255.255.0 broadcast 10.211.55.255
    inet6 fe80::21c:42ff:fea4:e10a prefixlen 64 scopeid 0x20<link>
    inet6 fdb2:2c26:f4e4:0:21c:42ff:fea4:e10a prefixlen 64 scopeid 0x0<global>
    ether 00:1c:42:a4:e1:0a txqueuelen 1000 (Ethernet)
    RX packets 385229 bytes 529902935 (505.3 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 133751 bytes 11330284 (10.8 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 1498 bytes 475914 (464.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1498 bytes 475914 (464.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:82:75:69 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "cd6dd4ff91585bca6e3253da53a8091f9a0fd2f0d2b1059d69c81129746ac7ac",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/cd6dd4ff9158",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {
    "none": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "442d24093083bb2d6bf5655b04dc0345ce2acdabb1e57c10b8476db5833cb705",
      "EndpointID": "6ac0d425fd76777ff981a47d5ea4282e9c7b34d3b84e0a75928cf6e2e00225a7",
      "Gateway": "",
      "IPAddress": "",
      "IPPrefixLen": 0,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "",
      "DriverOpts": null
    }
  }
}
```

PG容器

```
"NetworkSettings": {
  "Bridge": "",
  "SandboxID": "e3637cea3114abddc332931238bce563efd6aed1b8e4486367c26159191243d2",
  "HairpinMode": false,
  "LinkLocalIPv6Address": "",
  "LinkLocalIPv6PrefixLen": 0,
  "Ports": {},
  "SandboxKey": "/var/run/docker/netns/e3637cea3114",
  "SecondaryIPAddresses": null,
  "SecondaryIPv6Addresses": null,
  "EndpointID": "",
  "Gateway": "",
  "GlobalIPv6Address": "",
  "GlobalIPv6PrefixLen": 0,
  "IPAddress": "",
  "IPPrefixLen": 0,
  "IPv6Gateway": "",
  "MacAddress": "",
  "Networks": {
    "none": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "442d24093083bb2d6bf5655b04dc0345ce2acdabb1e57c10b8476db5833cb705",
      "EndpointID": "d28e263fb2df44c17498cfe3b540d2332d079063be84750b2a543d3932e17105",
      "Gateway": "",
      "IPAddress": "",
      "IPPrefixLen": 0,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "",
      "DriverOpts": null
    }
  }
}
```

网络模式-跨主机通信

Docker默认的网络环境下，单台主机的Docker容器可以通过docker0网桥直接通信，而不同主机上的Docker容器只能在宿主机上做端口映射进行通信。

这种端口映射方式对很多集群应用来说及不方便，如果能让Docker容器之间直接使用自己的IP地址通信，则会解决很多问题。

目前多主机通信方案按照原理可划分为：

- 直接路由方式；
- 桥接方式（如pipework）；
- Overlay隧道方式（如flannel、ovs+gre）；

网络模式-Pipework

Pipework是一个简易的Docker容器网络配置工具。由200多行shell脚本实现，可以通过使用ip、brctl、ovs-vsctl等命令来为Docker容器配置自定义的网桥、网卡、路由等。

```
sudo service docker stop
sudo ip link set dev docker0 down
sudo brctl delbr docker0
```

```
sudo brctl addbr br0
sudo ip addr add 10.211.55.1/24 dev br0
sudo ip link set dev br0 up
```

```
ip addr show br0
```

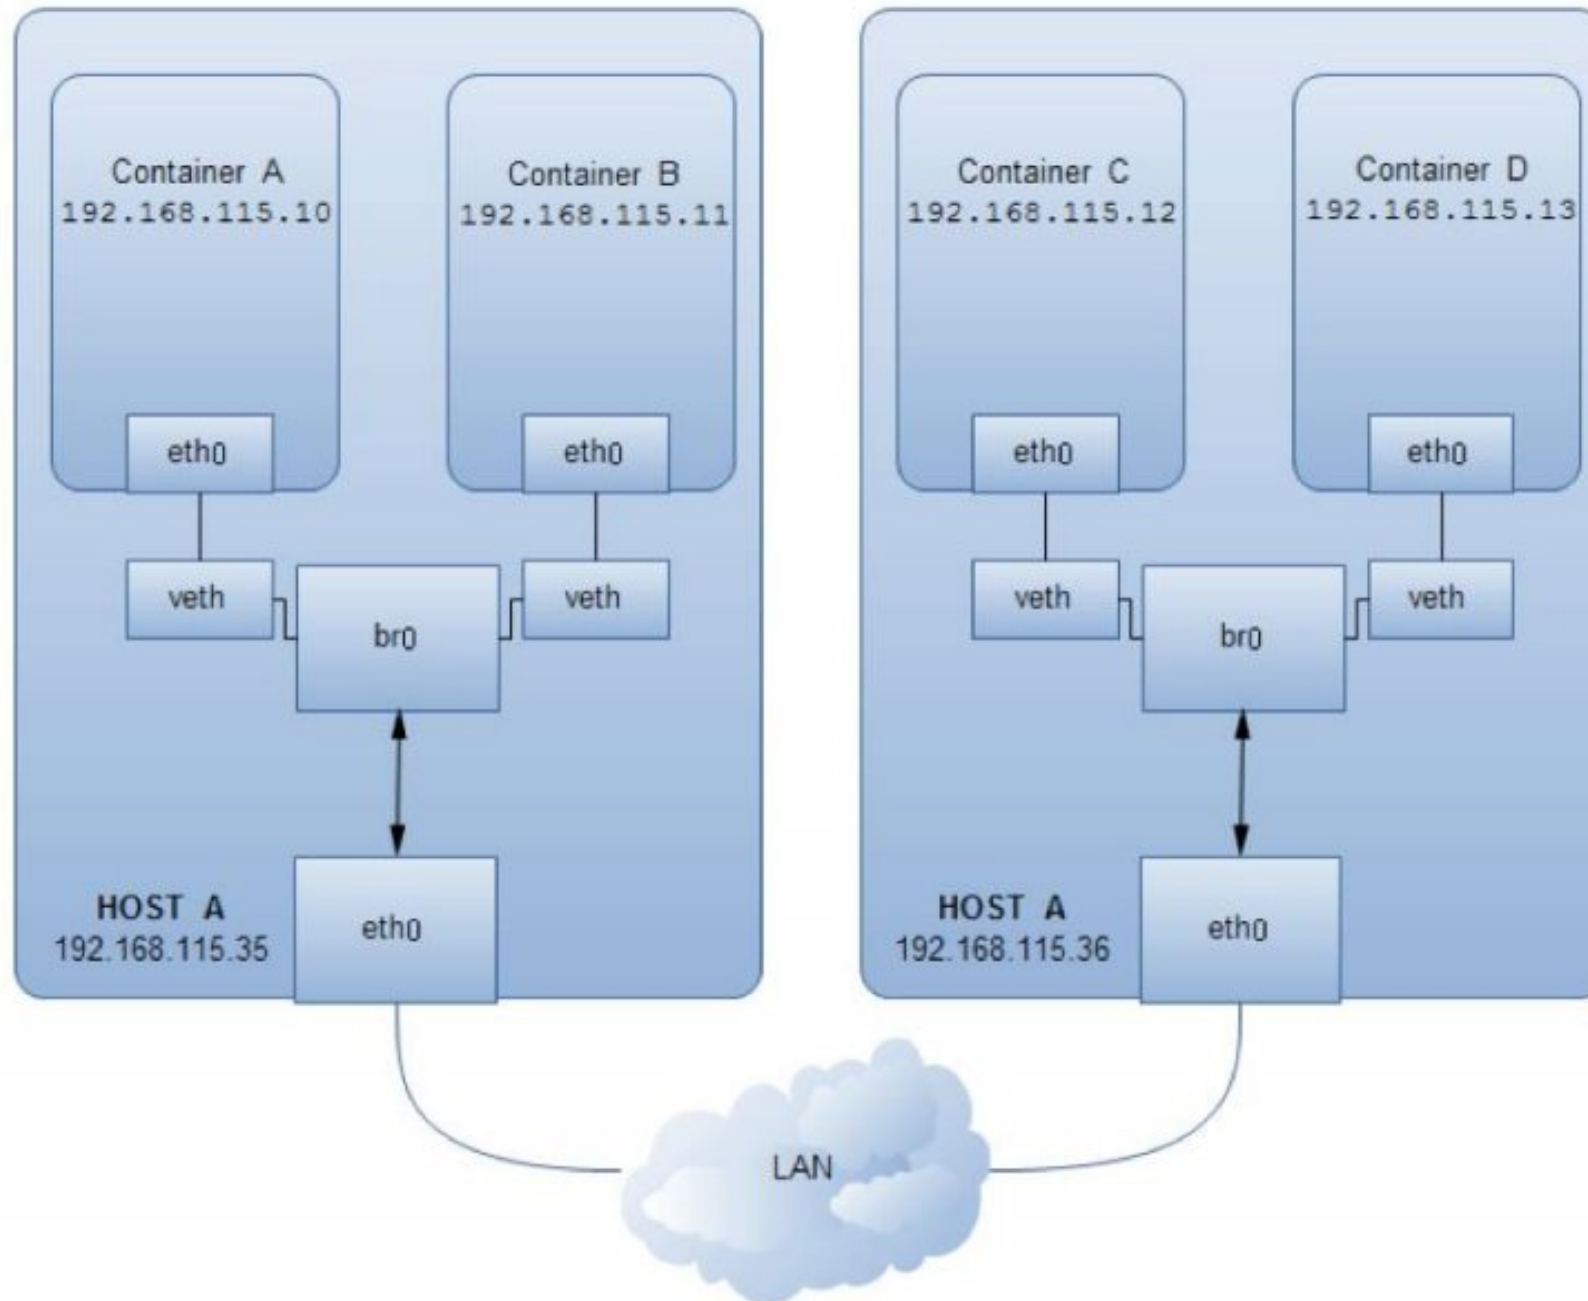
```
echo 'DOCKER_OPTS="-b=bridge0"' >> /etc/default/docker
```

```
sudo service docker start
```

```
[Service]
Type=notify
# the default is not to use systemd for c
# exists and systemd currently does not s
# for containers run by docker
ExecStart=/usr/bin/dockerd -b br0
ExecReload=/bin/kill -s HUP $MAINPID
```

网络模式-Pipework

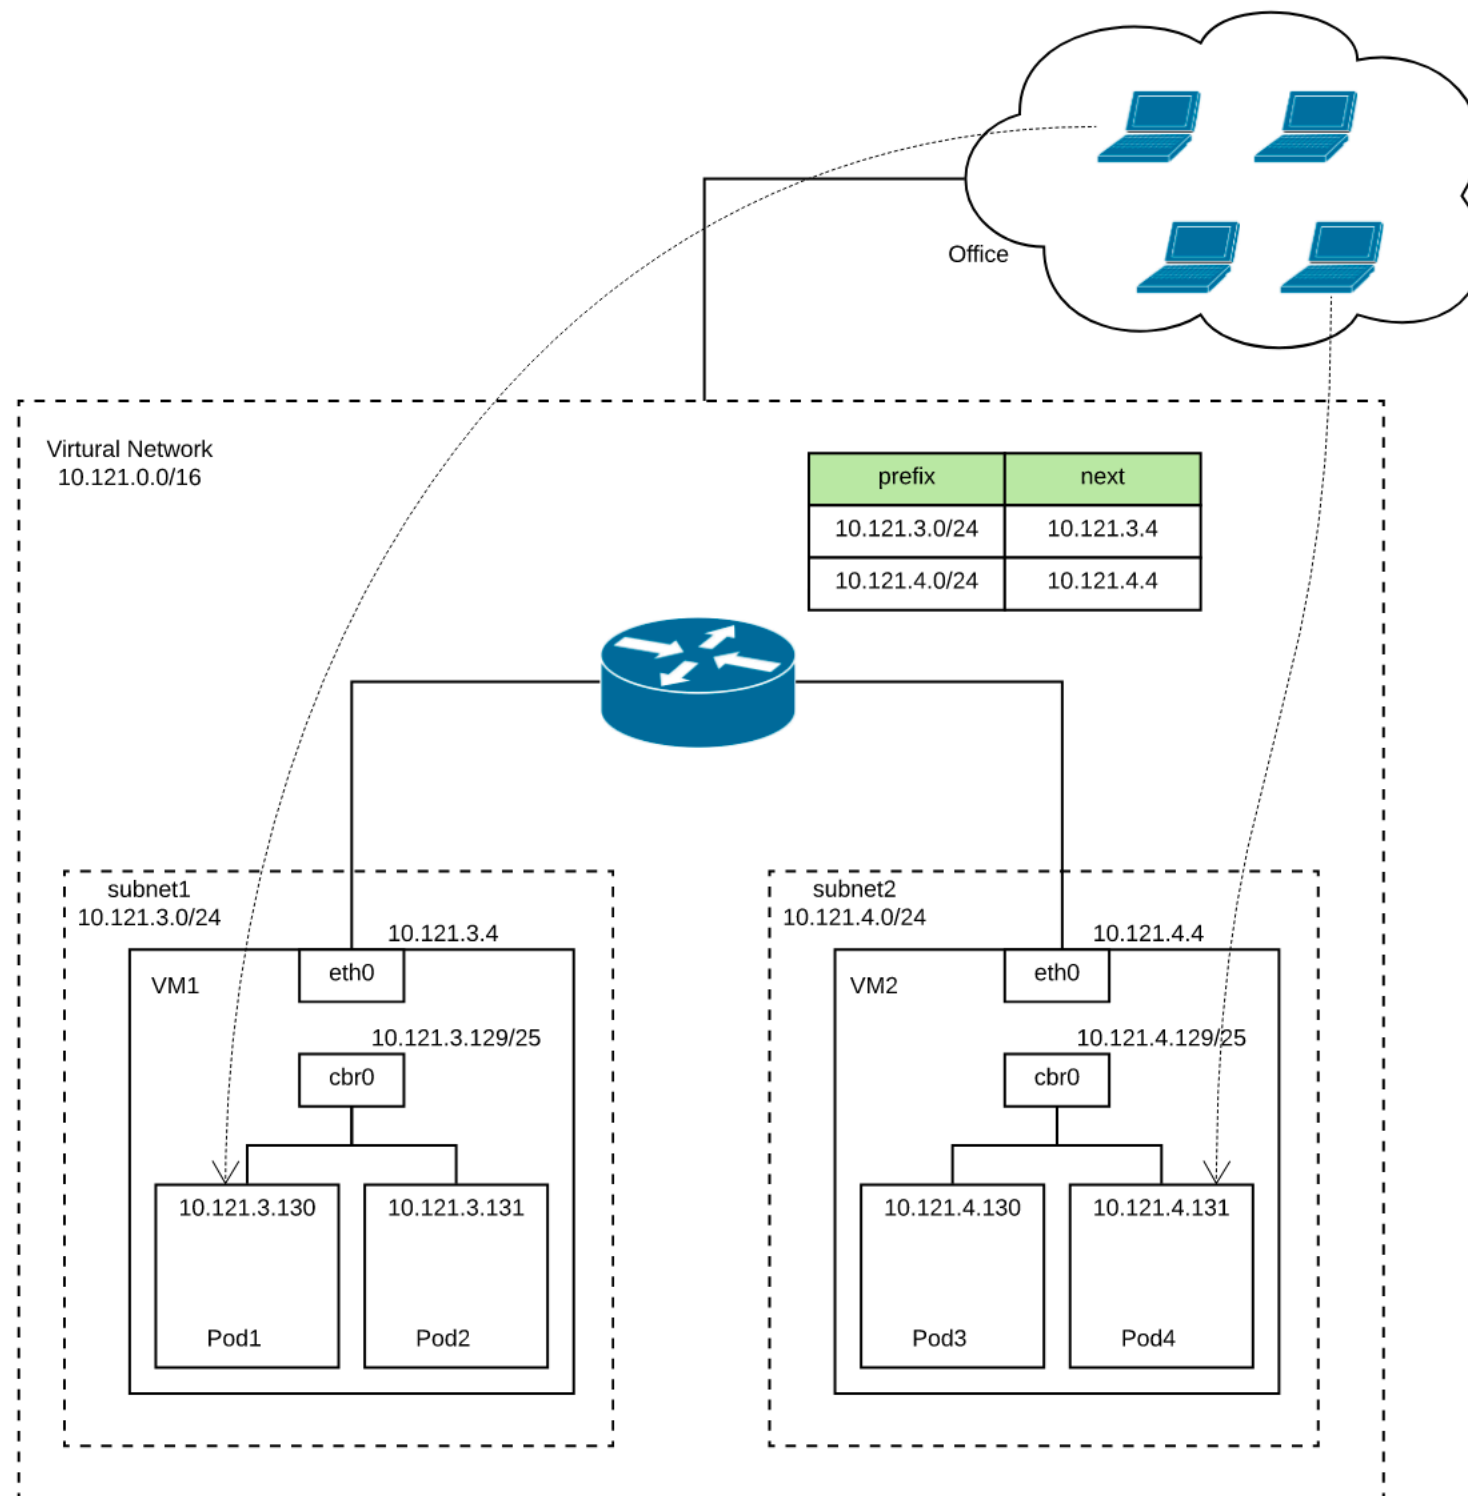
pipework



- 使用新建的`br0`网桥代替缺省的`docker0`网桥；
- `br0`网桥与`docker0`的区别在于`br0`和宿主机`eth0`之间是`veth pair`；

不同容器之间的通信可以借助于 pipework 这个工具给 docker 容器新建虚拟网卡并绑定 IP 桥接到 br0

Mobike实践



扁平网络，集群之间容器都属于相同网段，使得应用之间网络访问速度加快，同时减少管理负担

参考资料

<https://wiki.mobike.com/pages/viewpage.action?pageId=29278470>

Q&A